



Systems Software

Week 5: IPC and Pipes



Overview

- ↗ Inter Process Communication (IPC)
- ↗ File Descriptors
- ↗ Pipes
- ↗ Pipes Example
- ↗ Named Pipes (FIFO)

- ↗ Inter Process Communication (IPC)
- ↗ Facilitating communication between different processes.
- ↗ IPC can be offered between:
 - ↗ Related processes
 - ↗ Unrelated processes

File Streams and File Descriptors

- If we need to perform file operations (IO) to a file, there are 2 options available:
 - Streams
 - File Descriptors
- Streams are represented as `File *` objects.
- File descriptors are represented as objects of type `int`.

File Streams

- A stream offers a high level interface that is layered on top of the file descriptors.
- There are more features and functionality that comes with the stream interface for file IO.

File Descriptors

- File descriptors offer a primitive low level interface for IO operations.
- A file descriptor can connect to a file, a device (terminal), or a pipe or socket to communicate with another process.
- File descriptors should be used for IO with devices, and for nonblocking IO operations.

fopen and fdopen

- The **fopen()** function opens the file whose name is the string pointed to by path and associates a stream with it.
- The **fdopen()** function associates a stream with the existing file descriptor, fd. The mode of the stream (one of the values "r", "r+", "w", "w+", "a", "a+") must be compatible with the mode of the file descriptor.

fgetc

- `fgetc()` reads the next character from stream and returns it as an unsigned char cast to an int, or EOF on end of file or error.


fstat

- `stat()` retrieve information about the file pointed to by `pathname`; the differences for `fstatat()` are described below.
- `fstat()` is identical to `stat()`, except that the file about which information is to be retrieved is specified by the file descriptor `fd`.

What are Pipes

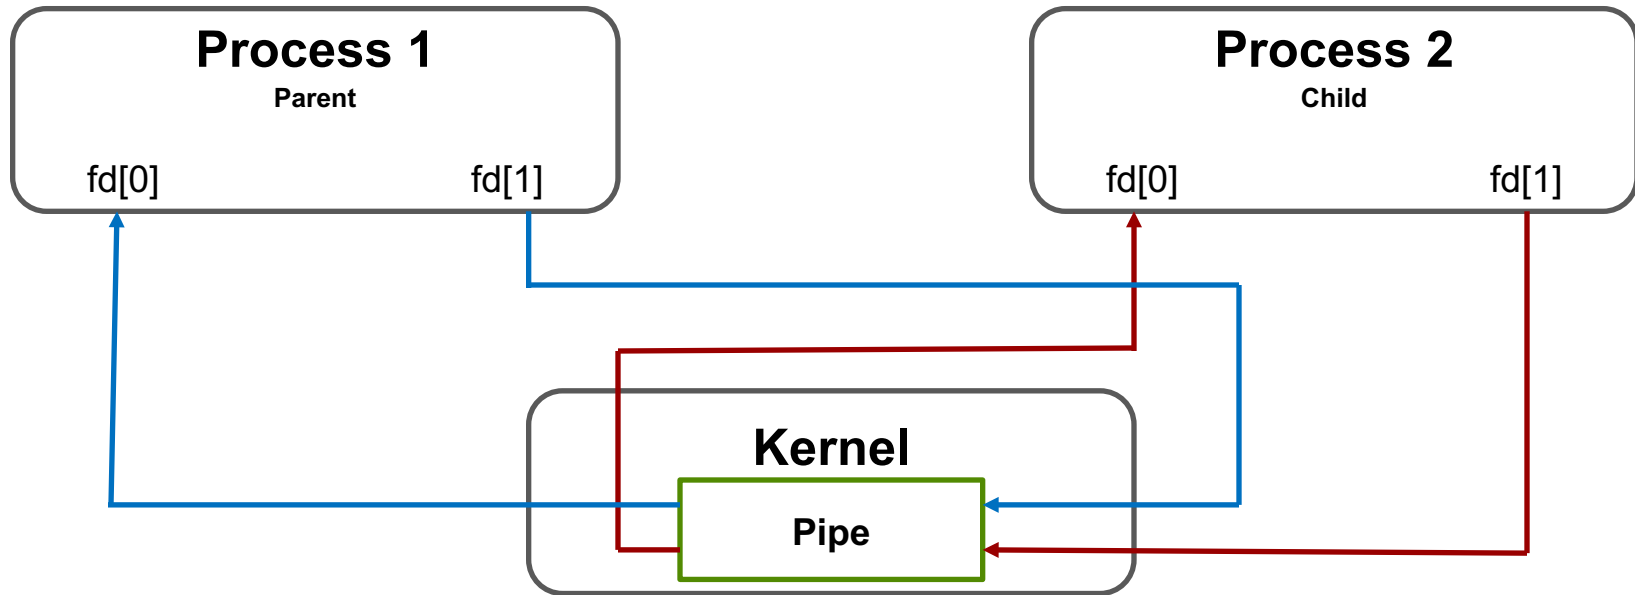
- A pipe allows processes to communicate with each other (Inter Process Communication - IPC)
- A process can write data to a pipe and this can be read by another process.
- The data passes to the pipe follows the FIFO algorithm.
- The pipe has no name.
- The pipe is created by a parent process and uses it to communicate with a child process.

Simple Pipe Architecture Example



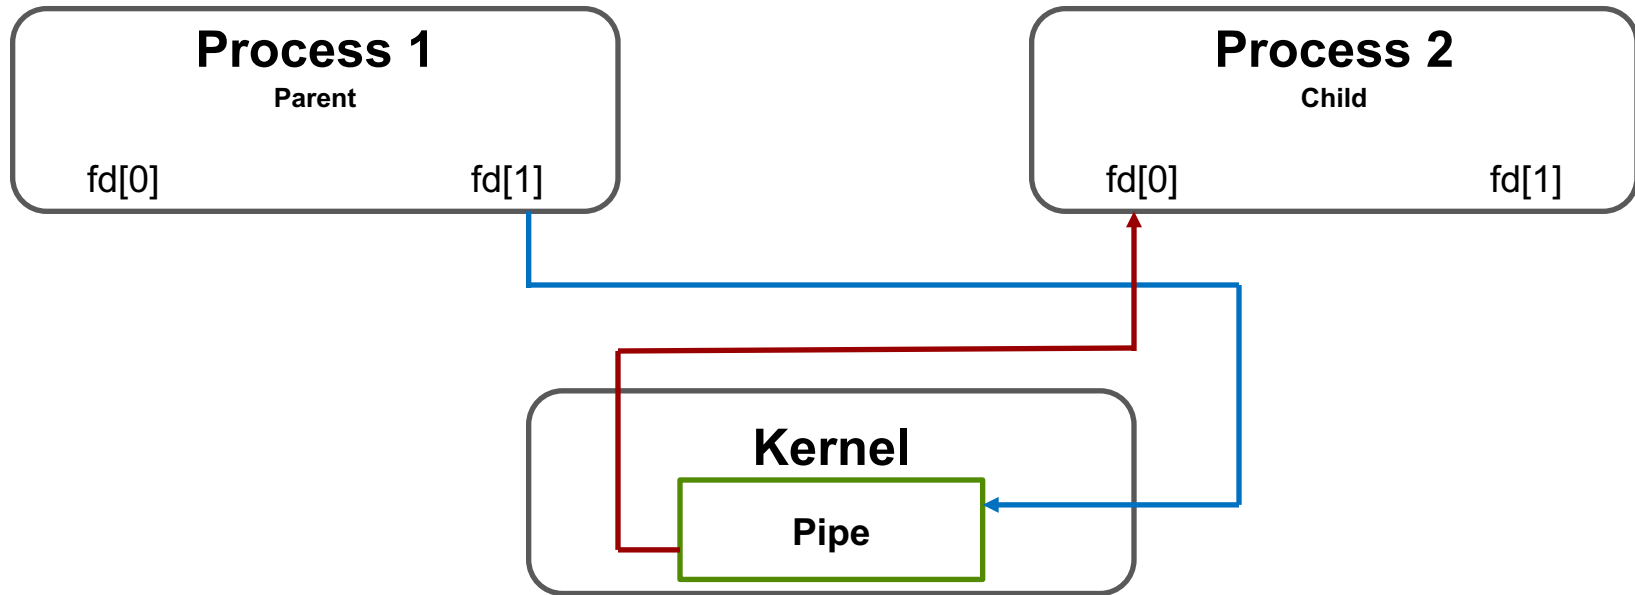
- Where fd = file descriptor
- fd[0] for reading
- fd[1] for writing
- The pipe is implemented using kernel memory.
- A pipe is a channel with two ends
- This example is a half-duplex pipe. Why??
- This example is of no real practical use, why would a single process need to communicate with itself?

Process forks



- After a fork we have to decide in which direction the data should flow
 - Parent to child
 - Child to parent

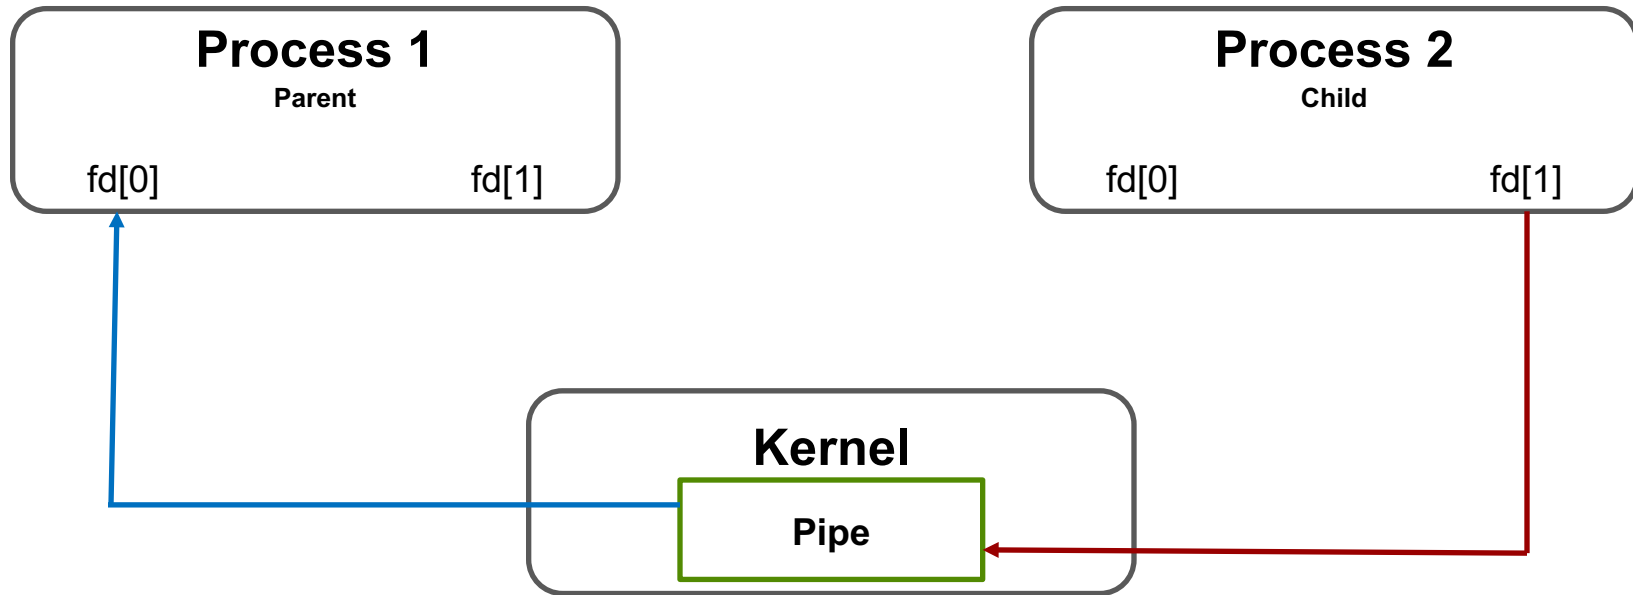
Parent to Child



➤ Parent must close `fd[0]`

➤ Child must close `fd[1]`

Child to Parent



➤ Parent must close `fd[1]`

➤ Child must close `fd[0]`

Rules for closing pipes::

- ↗ If we read from a pipe that was closed, the return should be 0
- ↗ If we write to a pipe with read closed on the other end, the SIGPIPE signal will be sent. This will need to be dealt with.

Creating a Pipe:: Step 1

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <stdlib.h>
#include <string.h>

int main(void)
{
```


Step 2:: Variables, pipe and fork

```
int      fd[2], nbytes;
pid_t    pid;
char     string[] = "Howya Friend!!!\n";
char     readbuffer[100];

// Create the pipe
pipe(fd);

// create a child process
pid = fork();
```

Step 3: Setup Child file Descriptors

```
if(pid == 0) // This is the child
{
    // Take no input, close fd[0]
    close(fd[0]);

    // Send output to parent via write
    write(fd[1], string, (strlen(string)+1));
    exit(0);
}
```

Step 4: Setup Child file Descriptors

```
else // This is the parent
{
    // Send no output, close fd[1]
    close(fd[1]);

    // Get input from the pipe via read
    nbytes = read(fd[0], readbuffer, sizeof(readbuffer));
    printf("Message from Child: %s", readbuffer);
}
```

Run

jmccarthy@debianJMC2017: ~/Documents/Apps/week5/pipes

x

File Edit View Search Terminal Help

```
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$ gcc -o pipe2 pipes2.c
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$ ./pipe2
Message from Child: Howya Friend!!!
jmccarthy@debianJMC2017:~/Documents/Apps/week5/pipes$
```

Pipe Example 2

- Create a C program to replicate the following:
 - `ps aux | grep login`
- How can pipes be used to solve this problem??

dup command

DUP(2)

Linux Programmer's Manual

DUP(2)

NAME

dup, dup2, dup3 - duplicate a file descriptor

SYNOPSIS

```
#include <unistd.h>
```

```
int dup(int oldfd);
```

```
int dup2(int oldfd, int newfd);
```

```
#define _GNU_SOURCE
```

```
#include <fcntl.h>
```

```
#include <unistd.h>
```

```
/* See feature_test_macros(7) */
```

```
/* Obtain 0_* constant definitions */
```

```
int dup3(int oldfd, int newfd, int flags);
```

dup command

`dup2()`

The **`dup2()`** system call performs the same task as **`dup()`**, but instead of using the lowest-numbered unused file descriptor, it uses the descriptor number specified in `newfd`. If the descriptor `newfd` was previously open, it is silently closed before being reused.

The steps of closing and reusing the file descriptor `newfd` are performed atomically. This is important, because trying to implement equivalent functionality using **`close(2)`** and **`dup()`** would be subject to race conditions, whereby `newfd` might be reused between the two steps. Such reuse could happen because the main program is interrupted by a signal handler that allocates a file descriptor, or because a parallel thread allocates a file descriptor.

Note the following points:

- * If `oldfd` is not a valid file descriptor, then the call fails, and `newfd` is not closed.
- * If `oldfd` is a valid file descriptor, and `newfd` has the same value as `oldfd`, then **`dup2()`** does nothing, and returns `newfd`.

dup command

`dup3()`

`dup3()` is the same as `dup2()`, except that:

- * The caller can force the close-on-exec flag to be set for the new file descriptor by specifying `O_CLOEXEC` in flags. See the description of the same flag in `open(2)` for reasons why this may be useful.
- * If oldfd equals newfd, then `dup3()` fails with the error `EINVAL`.

Pipes IPC Example 2

```
#include <stdlib.h>
#include <unistd.h>
```

```
void exec1();
void exec2();
```

```
int pid;
int pipefd[2];
```

- Init variables and functions.
- We have two functions that will use the exec command to perform different tasks in the process.
- Pid stores the process ID
- Pipefd array stores the file descriptors.

Pipes IPC Example 2

```
void main() {  
    // create pipe1  
    if (pipe(pipefd) == -1) {  
        perror("Error Init Pipe");  
        exit(1);  
    }  
}
```

- Start the program (main)
- Create a pipe and store the associated file descriptors in pipefd array

Pipes IPC Example 2

```
// fork (ps aux)
if ((pid = fork()) == -1) {
    perror("Error Init Fork 1");
    exit(1);
} else if (pid == 0) {
    exec1();
}
```

- Fork to create a child process
- If the pid is 0, we are dealing with the child process. Call the exec1 function to get the child process to swap to a different task.
- (exec1 function described on slide 27)

Pipes IPC Example 2

```
// fork (grep login)
if ((pid = fork()) == -1) {
    perror("Error Init Fork 1");
    exit(1);
} else if (pid == 0) {
    exec2();
}

close(pipefd[0]);
close(pipefd[1]);
} // close main
```

- Fork again to create a child process
- If the pid is 0, we are dealing with the child process. Call the exec2 function to get the child process to swap to a different task.
- (exec1 function described on slide 28)



Pipes IPC Example 2

```
void execl() {  
    // input from stdin - ok  
    // output to pipe  
    dup2(pipefd[1], 1);  
    // close fds  
    close(pipefd[0]);  
    close(pipefd[1]);  
    // exec  
    execlp("ps", "ps", "aux", NULL);  
    // exec didn't work, exit  
    perror("Error with ps aux");  
    exit(1);  
}
```

Pipes IPC Example 2

```
void exec2() {  
    // input from pipe  
    dup2(pipefd[0], 0);  
    // output to stdout - default  
  
    // close fds  
    close(pipefd[0]);  
    close(pipefd[1]);  
  
    // exec  
    execlp("grep", "grep", "login", NULL);  
    // exec didn't work, exit  
    perror("Error with grep login");  
    exit(1);  
}
```

Pipes – Making Life Easier!!

- The `popen()` function can be used to automate the creation of the pipe and the forking process.
- `fp = popen("ls *", "r");`
- `Popen` will set the file descriptors to facilitate the communication between parent and child or child and parent as required.

Pipes – Making Life Easier!!

popen.c x

*otherProcess.c x

```
#include <stdio.h>

int main() {
    FILE *fp;
    int status;
    char path[1024];

    fp = popen("ls *", "r");

    while (fgets(path, 1024, fp) != NULL)
        printf("%s", path);

    status = pclose(fp);
}
```

➤ Popen

1. Setup Pipes
2. Fork()
3. Child process runs `ls *` command and returns results to parent
4. Parent reads with `fgets` and displays with `printf`

Named Pipes

- A named pipe operates like a normal pipe
- The main difference is the named pipe exists as a special file
- The special file is a FIFO file, and doesn't contain any user info
- Processes of different ancestry can share data using named pipes.


FIFO file for a Named Pipe

- A named pipe operates like a file
- We can use the standard file IO system calls to operate a named pipe.
 - `int open(const char *pathname, int flags);`
 - `int read(int fd, void *buf, size_t count);`
 - `int write(int fd, const void *buf, size_t count);`
 - `int close(fd);`

FIFO Named Pipe using mkfifo

- The mkfifo function can be used to create a FIFO file.
- **mkfifo(fifoFile, 0666);**
- Use read and write operations to send information between processes.
- The processes don't need to be related (ancestors)
- Unlink can be called to remove the file.
- **unlink(fifoFile);**

FIFO Named Pipe Example



```
#include <fcntl.h>
#include <sys/stat.h>
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int fd;
    char * fifoFile = "/tmp/fifoFile";

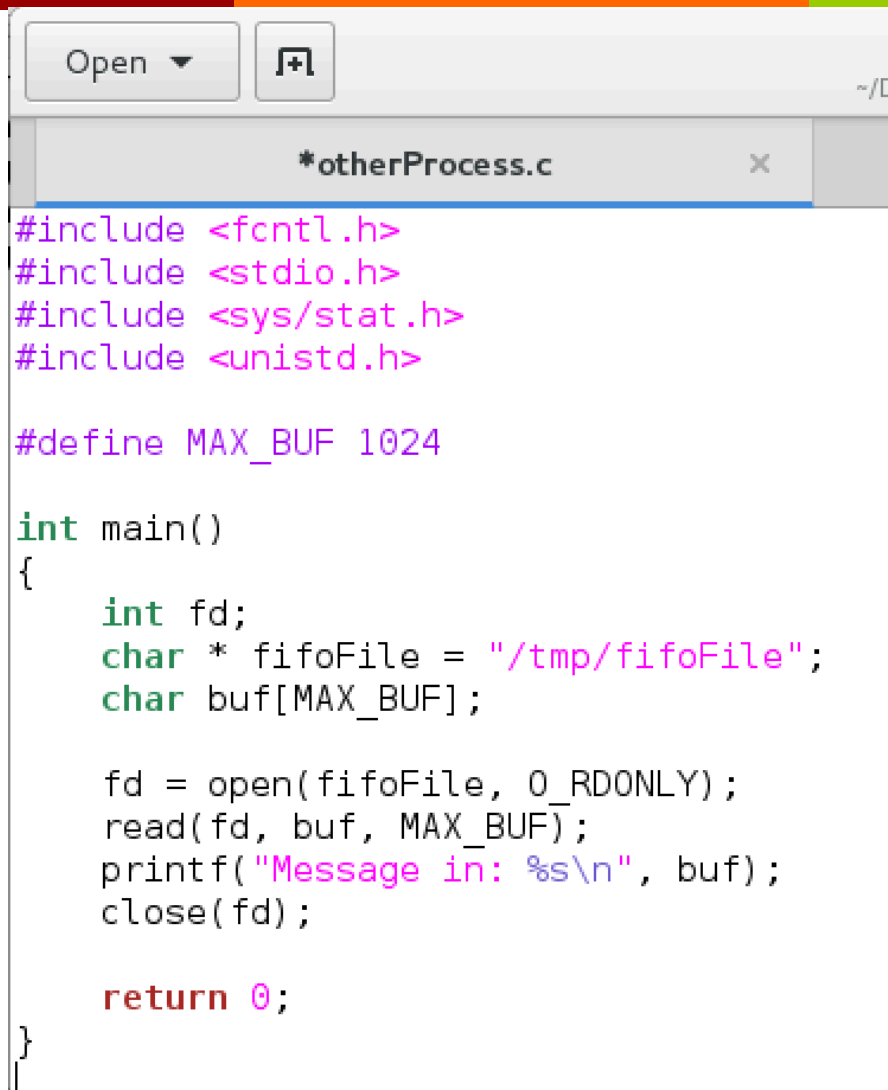
    /* create the FIFO (named pipe) */
    mkfifo(fifoFile, 0666);

    fd = open(fifoFile, O_WRONLY);
    write(fd, "The truth is out there!!", sizeof("The truth is out there!!"));
    close(fd);

    unlink(fifoFile);

    return 0;
}
```

FIFO Named Pipe Example



The image shows a code editor window with a title bar containing an "Open" button, a file icon, and a path indicator "~/.D". The editor has a tab labeled "*otherProcess.c" with a close button. The code is written in C and uses syntax highlighting: preprocessor directives and file paths are in magenta, keywords in green, and the return value in red.

```
#include <fcntl.h>
#include <stdio.h>
#include <sys/stat.h>
#include <unistd.h>

#define MAX_BUF 1024

int main()
{
    int fd;
    char * fifoFile = "/tmp/fifoFile";
    char buf[MAX_BUF];

    fd = open(fifoFile, O_RDONLY);
    read(fd, buf, MAX_BUF);
    printf("Message in: %s\n", buf);
    close(fd);

    return 0;
}
```

Run Example

```
jmccarthy@debianJMC2017: ~/Documents/A
File Edit View Search Terminal Help
$ ./mainProcess
$
```

```
jmccarthy@debianJMC2017: ~/Docu
File Edit View Search Terminal Help
$ ./otherProcess
Message in: The truth is out there!!
$
```

Questions

